

# Search Engine Prototype

Jaume Cloquell Capo

## Resumen

This project contain all the documentation of the search engine prototype.

## 1. Introduction

This project will use scikit-learn (Naive Bayes, MultinomialNB, CountVectorizer, TfidfTransformer, Bagging, Boosting) on Numpy and Pandas for the data frame;

## 2. Preprocess

- converting all letters to lower or upper case **DONE**
- converting numbers into words or removing numbers **TO DO**
- removing punctuations, accent marks and other diacritics **DONE**
- removing white spaces **DONE**
- expanding abbreviations **TO DO**
- removing stop words, sparse terms, and particular words **DOING**
  - Stemming: Stemming is a process of reducing words to their word stem, base or root form (for example, books—book, looked—look) **DONE**
  - Lemmatization: he aim of lemmatization, like stemming, is to reduce inflectional forms to a common base form. As opposed to stemming, lemmatization does not simply chop off inflections. Instead it uses lexical knowledge bases to get the correct base forms of words.**TO DO**
  - POS: Part-of-speech tagging aims to assign parts of speech to each word of a given text (such as nouns, verbs, adjectives, and others) based on its definition and its context. **TO DO**

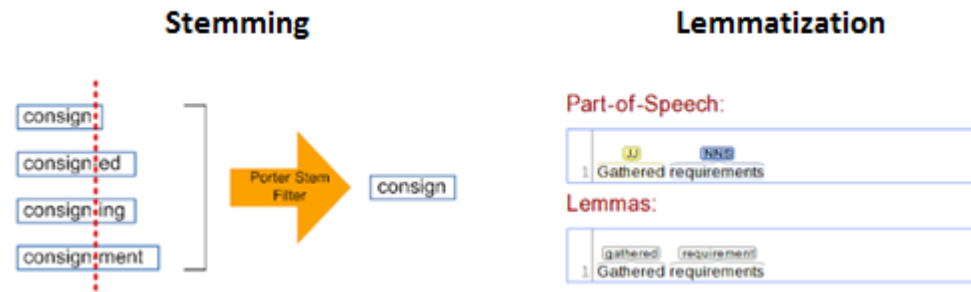


Figura 1: STEMMING VS Lemmatization

```

1000 #Filter only lang == EN
1001 data_en = data[data.locale == "en-US"]
1002 # Remove empty docs
1003 data_en = data_en[(data_en.medical_dictionary.notnull())]
1004 data = data_en[['id_specialization', 'medical_dictionary']]
1005 #data.medical_dictionary = data.medical_dictionary.apply(remove_html_tags)
1006 dictionary = {}
1007 for index, row in data.iterrows():
1008     # TODO: Remove words lower 1
1009     # Converting to Lowercase
1010     document = row.medical_dictionary
1011     document = remove_html_tags(document)
1012     document = document.lower()
1013     # Lemmatization
1014     document = document.split()
1015     document = [stemmer.lemmatize(word) for word in document]
1016     document = ' '.join(document)
1017     dictionary[row['id_specialization']] = [document]

```

### 3. Word2Vec and Doc2Ve

Google developed a method called Word2Vec that captures the context of words, while at the same time reducing the size of the data. Word2Vec is actually two different methods: Continuous Bag of Words (CBOW) and Skip-gram.

The above diagram is based on the CBOW model, but instead of using just nearby words to predict the word, we also added another feature vector, which is document-unique. So when training the word vectors  $W$ , the document vector  $D$  is trained as well, and in the end of training, it holds a numeric representation of the document.

The inputs consist of word vectors and document Id vectors. The word vector is a one-hot vector with a dimension  $1 \times V$ . The document Id vector has a dimension of  $1 \times C$ , where  $C$  is the number of total documents. The dimension of the weight matrix  $W$  of the hidden layer is  $V \times N$ . The dimension


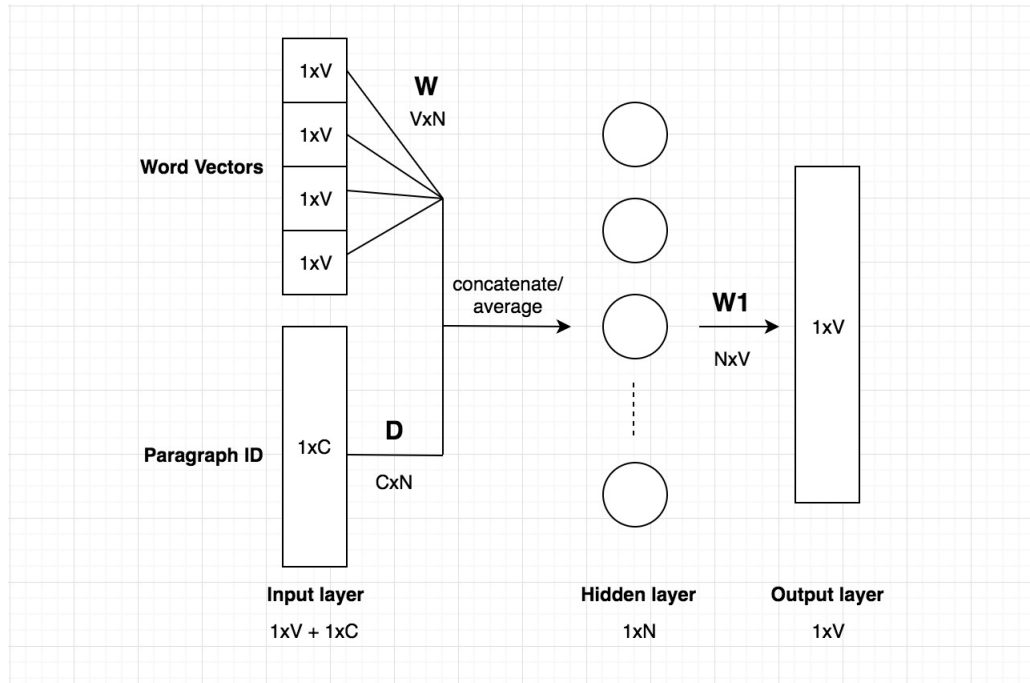


Figura 2: Doc2Vec

of the weight matrix  $D$  of the hidden layer is  $C \times N$ .

```
1000 train_documents = [TaggedDocument(data_en['tokens'][ind], str(data_en['id_specialization'][ind
1002   ]))
1003   for ind in data_en.index]
1004
1005   cores = multiprocessing.cpu_count()
1006   epochs = 1000
1007
1008   simple_models = [
1009       # PV-DM w/concatenation - window=5 (both sides) approximates paper's 10-word total
1010       # window size
1011       Doc2Vec(dm=1, dm_concat=1, vector_size=100, hs=0, min_count=2, workers=cores, epochs=
1012       epochs),
1013       # PV-DBOW
1014       Doc2Vec(dm=0, vector_size=100, hs=0, min_count=2, workers=cores, epochs=epochs),
1015       # PV-DM w/average
1016       Doc2Vec(dm=1, dm_mean=1, vector_size=100, hs=0, min_count=2, workers=cores, epochs=
1017       epochs),
1018   ]
1019
1020   model = simple_models[2]
1021
1022   model.build_vocab([x for x in train_documents])
1023
1024   model.train(train_documents, total_examples=model.corpus_count, epochs=model.epochs)
```

## 4. Models - Pros / Cons (Generalized)

### ■ KNN

#### ● Pros

- Is a non-parametric
- Variety of distance criteria to be choose from

#### ● Cons

- Define  $K$
- Slow algorithm
- Imbalanced data causes problems
- Outlier sensitivity

### ■ Naive Bayes

#### ● Pros

- Computationally fast

- Simple to implement
  - Works well with high dimensions
- Cons
  - Relies on independence assumption and will perform badly if this assumption is not met
- Random Forest
  - Pros
    - reduced variance
  - Cons
    - Not as easy to visually interpret
- Bagged Trees : train multiple trees using bootstrapped data to reduce variance and prevent overfitting
  - Pros
    - Reduces variance in comparison to regular decision trees
  - Cons
    - Does not reduce variance if the features are correlated
- Boosted Trees : Similar to bagging, but learns sequentially and builds off previous trees
  - Pros
    - Somewhat more interpretable than bagged trees/random forest as the user can define the size of each tree resulting in a collection of stumps (1 level) which can be viewed as an additive model
  - Cons
    - Unlike bagging and random forests, can overfit if number of trees is too large

## 5. Code

$$\mathbf{tf}(t, d) = \frac{f_d(t)}{\max_{w \in d} f_d(w)}$$

$$\mathbf{idf}(t, D) = \ln \left( \frac{|D|}{|\{d \in D : t \in d\}|} \right)$$

$$\mathbf{tfidf}(t, d, D) = \mathbf{tf}(t, d) \cdot \mathbf{idf}(t, D)$$

$$\mathbf{tfidf}'(t, d, D) = \frac{\mathbf{idf}(t, D)}{|D|} + \mathbf{tfidf}(t, d, D)$$

$f_d(t) :=$  frequency of term  $t$  in document  $d$

$D :=$  corpus of documents

Figura 3: Tf idf

```

1000 trclass SearchEngine():
1002     def __init__(self, label_names, X_train, y_train):
1004         self.k = len(y_train) # K is the number of classes, in this case, specializations
1006         self.label_names = label_names
1008         self.X_train, self.y_train = X_train, y_train
1010
1012     def fit(self):
1014         # min_df: This corresponds to the minimum number of documents that should contain this
1016         # feature.
1018         # max_df: we should include only those words that occur in a maximum of 70% of all the
1020         # documents
1022         self.vectorizer = CountVectorizer(ngram_range=(1, 1), max_features=1500, min_df=5,
1024         max_df=0.4, stop_words=stopwords.words('english'))
1026
1028         X_train_vect = self.vectorizer.fit_transform(self.X_train)
1030         self.tfidf_transformer = TfidfTransformer()
1032         X_train_trans = self.tfidf_transformer.fit_transform(X_train_vect)
1034
1036         #self.classifier = KNeighborsClassifier(n_neighbors=self.k)
1038         #self.classifier = RandomForestClassifier(n_estimators=500, max_features=0.25,
1040         criterion="entropy", class_weight="balanced")
1042         #self.classifier = BaggingClassifier(n_estimators =25, max_features=0.25)
1044         #self.classifier = GradientBoostingClassifier(n_estimators =100, learning_rate =0.1,
1046         max_depth=6, min_samples_leaf =1, max_features=1.0) clf.fit(X, training_set_y)
1048         self.classifier = MultinomialNB()
1050
1052         self.classifier.fit(X_train_trans, self.y_train)
1054
1056     def predict(self, X_test):
1058         X_test_vect = self.vectorizer.transform(X_test)
1060         X_test_trans = self.tfidf_transformer.transform(X_test_vect)
1062         y_pred = self.classifier.predict(X_test_trans)
1064         return y_pred

```

```

1030 def predict_single(self, doc):
1031     X_test_vect = self.vectorizer.transform([doc])
1032     X_test_trans = self.tfidf_transformer.transform(X_test_vect)
1033     y_pred = zip(self.classifier.classes_, self.classifier.predict_proba(X_test_trans)[0])
1034     y_pred = sorted([(self.label_names[ind], score) for ind, score in y_pred], key=lambda
x: -x[1])
1035     return y_pred
1036
1037 def report(self, X_test, y_test, y_pred):
1038     print(classification_report(y_test, y_pred, target_names=self.label_names, digits=4))
1039
1040     total = 0
1041     same = 0
1042     for i in range(len(y_test)):
1043         if y_test[i] == y_pred[i]:
1044             same += 1
1045             total += 1
1046     print(total, same)

```

## 6. What I would do

- MultiLanguage
  - MultiLanguage embeddings <https://github.com/facebookresearch/MUSE>
  - Diferent modals
- Add articles to the dataset
- Visit
  - ElasticSearch: <https://www.elastic.co/es/>
  - Spacy <https://spacy.io/>
  - Keras workshop <https://github.com/tensorflow/workshops/blob/master/extras/keras-bag-of-keras-bow-model.ipynb>
- POS
- $\text{TF-IDF}(\text{document}) = \text{TF-IDF}(\text{title}) * \alpha + \text{TF-IDF}(\text{body}) * (1-\alpha)$ 
  - Calculate TF-IDF for Body for all docs
  - Calculate TF-IDF for title for all docs
  - Multiply the Body TF-IDF with alpha